


# IF3140 – Sistem Basis Data Concurrency Control:

- Introduction
- Lock-based Protocol




KNOWLEDGE & SOFTWARE ENGINEERING





KNOWLEDGE & SOFTWARE ENGINEERING

©Silberschatz et.al. (2019)  
[modified by Tim Pengajar IF3140 Semester 1 2024/2025]



# Sumber

Silberschatz, Korth, Sudarshan:  
“Database System Concepts”,  
7<sup>th</sup> Edition

- Chapter 18:  
Concurrency Control

# *Objectives*

Students are able to:

- Explain the effect of different isolation levels on the concurrency control mechanisms
- Choose the proper isolation level for implementing a specified transaction protocol



KNOWLEDGE & SOFTWARE ENGINEERING

# Outline



KNOWLEDGE & SOFTWARE ENGINEERING

- Lock-Based Protocols
  - 2-phase locking
  - Graph-based protocols
  - Deadlock handling
  - Multiple granularity
- Timestamp-Based Protocols
- Validation-Based Protocols
- Insert and delete operations
- Multiversion Schemes
  - MV Timestamp ordering
  - MV 2-phase locking
  - Snapshot isolation
- Weak levels of consistency

# Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
  1. **exclusive** (X) mode. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
  2. **shared** (S) mode. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.

# *Lock-Based Protocols (Cont.)*

- **Lock-compatibility matrix**

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item,
- But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.



KNOWLEDGE & SOFTWARE ENGINEERING


# Schedule With Lock Grants

- Grants omitted in rest of chapter
  - Assume grant happens just before the next instruction following lock request
- This schedule is not serializable (why?)
- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols enforce serializability by restricting the set of possible schedules.

$T_1$	$T_2$	concurrency-control manager
lock-X( $B$ )		grant-X( $B, T_1$ )
read( $B$ )		
$B := B - 50$		
write( $B$ )		
unlock( $B$ )		
	lock-S( $A$ )	grant-S( $A, T_2$ )
	read( $A$ )	
	unlock( $A$ )	
	lock-S( $B$ )	grant-S( $B, T_2$ )
	read( $B$ )	
	unlock( $B$ )	
	display( $A + B$ )	
lock-X( $A$ )		grant-X( $A, T_1$ )
read( $A$ )		
$A := A + 50$		
write( $A$ )		
unlock( $A$ )		

# The Two-Phase Locking Protocol

- A protocol which ensures conflict-serializable schedules.
- Phase 1: **Growing Phase**
  - Transaction may obtain locks
  - Transaction may not release locks
- Phase 2: **Shrinking Phase**
  - Transaction may release locks
  - Transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock).



# Deadlock

- Consider the partial schedule

$T_3$	$T_4$
lock-X( $B$ ) read( $B$ ) $B := B - 50$ write( $B$ )	
lock-X( $A$ )	lock-S( $A$ ) read( $A$ ) lock-S( $B$ )

- Neither  $T_3$  nor  $T_4$  can make progress — executing **lock-S( $B$ )** causes  $T_4$  to wait for  $T_3$  to release its lock on  $B$ , while executing **lock-X( $A$ )** causes  $T_3$  to wait for  $T_4$  to release its lock on  $A$ .
- Such a situation is called a **deadlock**.
  - To handle a deadlock one of  $T_3$  or  $T_4$  must be rolled back and its locks released.



KNOWLEDGE & SOFTWARE ENGINEERING

# Deadlock (Cont.)

- The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.
- **Starvation** is also possible if concurrency control manager is badly designed. For example:
  - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
  - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.

# *The Two-Phase Locking Protocol (Cont.)*

- Two-phase locking does not ensure freedom from deadlocks
- Extensions to basic two-phase locking needed to ensure recoverability or freedom from cascading roll-back
  - **Strict two-phase locking:** a transaction must hold all its exclusive locks till it commits/aborts.
    - Ensures recoverability and avoids cascading roll-backs
  - **Rigorous two-phase locking:** a transaction must hold *all* locks till commit/abort.
    - Transactions can be serialized in the order in which they commit.
- Most databases implement rigorous two-phase locking, but refer to it as simply two-phase locking

# The Two-Phase Locking Protocol (Cont.)

- Two-phase locking is not a necessary condition for serializability
  - There are conflict serializable schedules that cannot be obtained if the two-phase locking protocol is used.
- In the absence of extra information (e.g., ordering of access to data), two-phase locking is necessary for conflict serializability *in the following sense*:
  - Given a transaction  $T_i$  that does not follow two-phase locking, we can find a transaction  $T_j$  that uses two-phase locking, and a schedule for  $T_i$  and  $T_j$  that is not conflict serializable.

$T_1$	$T_2$
lock-X( $B$ )	
read( $B$ )	
$B := B - 50$	
write( $B$ )	
unlock( $B$ )	
	lock-S( $A$ )
	read( $A$ )
	unlock( $A$ )
	lock-S( $B$ )
	read( $B$ )
	unlock( $B$ )
	display( $A + B$ )
	lock-X( $A$ )
	read( $A$ )
	$A := A + 50$
	write( $A$ )
	unlock( $A$ )

# *Locking Protocols*

- Given a locking protocol (such as 2PL)
  - A schedule  $S$  is **legal** under a locking protocol if it can be generated by a set of transactions that follow the protocol
  - A protocol **ensures** serializability if all legal schedules under that protocol are serializable



KNOWLEDGE & SOFTWARE ENGINEERING

# Exercise 9.1

Consider the following two transactions:

**T1:** **read (A)**  
**read (B)**  
**B := B + 0.1\*A**  
**write (B)**

**T2:** **read (B)**  
**read (A)**  
**A := A - 0.05\*B**  
**write (A)**

- Add lock and unlock instructions to both transactions so that they follow the two-phase locking protocol.
- Can the execution of these two transactions result in a deadlock?

# *Exercise 9.2*

Is the following schedule a two-phase locking (2PL) schedule (legal under 2PL protocol)?

**R1(A); R2(A); R3(B); W1(A); R2(C); R2(B); W2(B); W1(C);**



KNOWLEDGE & SOFTWARE ENGINEERING

**Solution**

©Silberschatz et.al. (2019)

[modified by Tim Pengajar IF3140 Semester 1 2024/2025]

# Lock Conversions

- Two-phase locking protocol with lock conversions:
  - Growing Phase:
    - can acquire a lock-S on item
    - can acquire a lock-X on item
    - can **convert** a lock-S to a lock-X (**upgrade**)
  - Shrinking Phase:
    - can release a lock-S
    - can release a lock-X
    - can convert a lock-X to a lock-S (**downgrade**)
- This protocol ensures serializability

# *Automatic Acquisition of Locks*

- A transaction  $T_i$  issues the standard read/write instruction, without explicit locking calls.
- The operation **read**( $D$ ) is processed as:

**if**  $T_i$  has a lock on  $D$

**then**

read( $D$ )

**else begin**

if necessary wait until no other  
transaction has a **lock-X** on  $D$

grant  $T_i$  a **lock-S** on  $D$ ;

read( $D$ )

**end**



KNOWLEDGE & SOFTWARE ENGINEERING

# *Automatic Acquisition of Locks (Cont.)*

- The operation **write**(D) is processed as:

```
if  $T_i$  has a lock-X on D  
then  
    write(D)  
else begin  
    if necessary wait until no other trans. has any lock on D,  
    if  $T_i$  has a lock-S on D  
        then  
            upgrade lock on D to lock-X  
        else  
            grant  $T_i$  a lock-X on D  
    write(D)  
end;
```

- All locks are released after commit or abort

# Exercise 9.3

Instructions from T1, T2, and T3 arrive in the following order.

**R1(A); R2(A); R3(B); W1(A); R2(C); R2(B); C3; W2(B); C2;  
W1(C); C1;**

What is the final schedule if the 2-phase locking with automatic acquisition of locks is implemented by CC Manager?