



# IF3170 Artificial Intelligence Regression



Masayu Leylia Khodra

# Capaian Pembelajaran

- Mahasiswa dapat menjelaskan dan menerapkan penggunaan konsep supervised learning
- Mahasiswa dapat menjelaskan dan menerapkan teknik linear regression untuk persoalan pembelajaran mesin yang sesuai.

# Regression: Introduction

As an example, we are a real estate agent, and in charge of selling a new house. We don't know the price, and we want to infer it by comparing it with other houses. We look at features of the house which could influence the house, i.e. number of rooms. At the end of the day, what we want is **a formula on the feature** which gives us the price of the house, or at least an estimate for it.



# Linear Relationship

- The price of the house is **dependent variable** or **response**, and number of rooms is **independent variable** or **regressor** or **predictor**.
- A reasonable form of a relationship between the response Y and the regressor x is the linear relationship:  $Y = \beta_0 + \beta_1 x$






Figure 11.1: A linear relationship;  $\beta_0$ : intercept;  $\beta_1$ : slope.

# Regression: Predict Response Y based on Regressor x

Figure 6. Our task is now to predict the price of the house with 4 rooms. Using the model (line), we deduce that the predicted price of this house is \$300.



# Regression Modeling: Example




```
1 df.shape
```

(1067, 13)

```
1 df[['ENGINESIZE', 'CO2EMISSIONS']].sample(5)
```

	ENGINESIZE	CO2EMISSIONS
697	2.0	196
872	3.8	253
614	2.4	200
351	5.4	382
426	3.5	258




# Simple Linear Regression: Training

```
1 from sklearn import linear_model  
2 regr = linear_model.LinearRegression()  
3 train_x = np.asarray(train[['ENGINESIZE']])  
4 train_y = np.asarray(train[['CO2EMISSIONS']])  
5 regr.fit (train_x, train_y)  
6 # The coefficients  
7 print ('Coefficients: ', regr.coef_)  
8 print ('Intercept: ',regr.intercept_)
```

Coefficients: [[39.43522758]]

Intercept: [124.41641136]




$$y=125.49+39.14x$$

# Multivariate Linear Regression: Training & Testing

```
1 from sklearn import linear_model  
2 regr2 = linear_model.LinearRegression()  
3 train_x = np.asarray(train[['ENGINESIZE','CYLINDERS','FUELCONSUMPTION_COMB']])  
4 train_y = np.asarray(train[['CO2EMISSIONS']])  
5 regr2.fit (train_x, train_y)  
6 print ('Coefficients: ', regr2.coef_)  
7 print ('Intercept: ',regr2.intercept_)
```

Coefficients: [[11.36694334 7.26349823 9.58219844]]  
Intercept: [64.98474043]

	ENGINESIZE	CYLINDERS	FUELCONSUMPTION_COMB	CO2EMISSIONS
0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136
3	3.5	6	11.1	255
4	3.5	6	10.6	244



# Estimating Parameters: Least Square Estimator

- Objective function: minimize sum of the squares of the residuals/errors (SSE)

$$SSE = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - b_0 - b_1 x_i)^2$$

$$\frac{\partial(SSE)}{\partial b_0} = -2 \sum_{i=1}^n (y_i - b_0 - b_1 x_i)$$

$$nb_0 + b_1 \sum_{i=1}^n x_i = \sum_{i=1}^n y_i$$

$$b_0 = \frac{\sum_{i=1}^n y_i - b_1 \sum_{i=1}^n x_i}{n} = \bar{y} - b_1 \bar{x}$$

$$\frac{\partial(SSE)}{\partial b_1} = -2 \sum_{i=1}^n (y_i - b_0 - b_1 x_i) x_i$$

$$b_1 = \frac{n \sum_{i=1}^n x_i y_i - \left( \sum_{i=1}^n x_i \right) \left( \sum_{i=1}^n y_i \right)}{n \sum_{i=1}^n x_i^2 - \left( \sum_{i=1}^n x_i \right)^2} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \sigma_{XY}/\sigma_x^2$$

Let  $X$  be a random variable with probability distribution  $f(x)$  and mean  $\mu$ . The variance of  $X$  is

$$\sigma^2 = E[(X - \mu)^2] = \sum_x (x - \mu)^2 f(x), \quad \text{if } X \text{ is discrete, and}$$

$$\sigma^2 = E[(X - \mu)^2] = \int_{-\infty}^{\infty} (x - \mu)^2 f(x) dx, \quad \text{if } X \text{ is continuous.}$$

The positive square root of the variance,  $\sigma$ , is called the **standard deviation** of  $X$ .

Let  $X$  and  $Y$  be random variables with joint probability distribution  $f(x, y)$ . The covariance of  $X$  and  $Y$  is

$$\sigma_{XY} = E[(X - \mu_X)(Y - \mu_Y)] = \sum_x \sum_y (x - \mu_X)(y - \mu_Y) f(x, y)$$

if  $X$  and  $Y$  are discrete, and

$$\sigma_{XY} = E[(X - \mu_X)(Y - \mu_Y)] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} (x - \mu_X)(y - \mu_Y) f(x, y) dx dy$$

if  $X$  and  $Y$  are continuous.

Let  $X$  and  $Y$  be random variables with covariance  $\sigma_{XY}$  and standard deviations  $\sigma_X$  and  $\sigma_Y$ , respectively. The correlation coefficient of  $X$  and  $Y$  is

$$\rho_{XY} = \frac{\sigma_{XY}}{\sigma_X \sigma_Y}.$$

# Implement From Scratch: Simple Linear Regression

## From Scratch

```
b1=covariance(X,Y)/variance(X)  
b0=mean(Y)-b1*mean(X)
```

```
1 # Coefficients: [[39.43522758]]  
2 # Intercept: [124.41641136]  
3 print('cov(X,Y)=',train.CO2EMISSIONS.cov(train.ENGINESIZE))  
4 print('var(X)=',train.ENGINESIZE.var())  
5 b1=train.CO2EMISSIONS.cov(train.ENGINESIZE)/train.ENGINESIZE.var()  
6 b0=train.CO2EMISSIONS.mean()-b1*train.ENGINESIZE.mean()  
7 print('Coefficient=',b1)  
8 print('Intercept=',b0)
```






```
cov(X,Y)= 78.95604822766053  
var(X)= 2.0021704722919877  
Coefficient= 39.43522757943557  
Intercept= 124.41641136306387
```

```
1 from sklearn import linear_model  
2 regr = linear_model.LinearRegression()  
3 train_x = np.asarray(train[['ENGINESIZE']])  
4 train_y = np.asarray(train[['CO2EMISSIONS']])  
5 regr.fit (train_x, train_y)  
6 # The coefficients  
7 print ('Coefficients: ', regr.coef_)  
8 print ('Intercept: ',regr.intercept_)
```

```
Coefficients: [[39.43522758]]  
Intercept: [124.41641136]
```



# Linear Regression: Supervised Learning

Learning a (possibly incorrect) general function from specific input-output pairs



$$y = 64.98 + 11.37x_1 + 7.26x_2 + 9.58x_3$$

	ENGINESIZE	CYLINDERS	FUELCONSUMPTION_COMB	CO2EMISSIONS
0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136
3	3.5	6	11.1	255
4	3.5	6	10.6	244



	ENGINESIZE	CYLINDERS	FUELCONSUMPTION_COMB
	2.0	4	8.5
	2.4	4	9.6
	1.5	4	5.9

```
1 test_y[0:3]
```

```
array([[196],  
       [221],  
       [136]], )
```

$$y = 64.98 + 11.37x_1 + 7.26x_2 + 9.58x_3$$

```
1 #test_y_ = regr2.predict(test_x)  
2 test_y_[0:3]  
array([[198.22130673],  
       [213.30850235],  
       [167.62411912]])
```

# Linear Regression: Testing

```
1 from sklearn.metrics import r2_score
2
3 test_x = np.asarray(test[['ENGINESIZE']])
4 test_y = np.asarray(test[['CO2EMISSIONS']])
5 test_y_ = regr.predict(test_x)
6
7 print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
8 print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
9 print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```


Mean absolute error: 22.50  
Residual sum of squares (MSE): 840.97  
R2-score: 0.73

$$y=125.49+39.14x$$

$$y=64.98+11.37x_1+7.26x_2+9.58x_3$$

```
1 from sklearn.metrics import r2_score
2
3 test_x = np.asarray(test[['ENGINESIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB']])
4 test_y_ = regr2.predict(test_x)
5
6 print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_ - test_y)))
7 print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_ - test_y) ** 2))
8 print("R2-score: %.2f" % r2_score(test_y_ , test_y) )
```

Mean absolute error: 15.20  
Residual sum of squares (MSE): 443.90  
R2-score: 0.87



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Coefficient of Determination, R<sup>2</sup>

$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2},$$

# Linear Regression - LSE: Summary

**Target function** is estimated by **hypothesis** as a linear function of X:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \dots + \theta_n x_n = \sum \theta_i x_i, i \text{ in } [0..n]$$
$$\rightarrow h_{\theta}(x) = \theta^T \cdot x$$

**Cost function**  $J(\theta)$  assigns a large cost to bad predictions  $h(x)$  and small cost to good predictions  $h(x)$ . For LSE, we used **sum of squares errors (SSE)**.

Find good  $\theta$  (parameters or weights):  $\theta^* = \operatorname{argmin} J(\theta)$

Analytical solution: find  $\theta$  that solve  $\nabla J(\theta) = 0$ .

If  $\nabla J(\theta) = 0$  is unable to solve, use iterative numerical methods. The simplest method is gradient descent.

# Pertanyaan ?